



Why Model-Level Offloading Fails: Transfer Dominates Computation


(a) Model-Level Offloading: Sequential Transfer → Compute



(b) Block-Level Sequential: No Pipelining (Transfer→Compute per block)



(c) Block-Level Pipelined (RabbitVideo): Overlap Transfer[i+1] with Compute[i]



PCIe Transfer (CPU→GPU) GPU Computation